

UPPSALA UNIVERSITET



HIGH PERFORMANCE PROGRAMMING 1TD062

ASSIGNMENT 3 - GALAXY SIMULATION

Agelii, Carl.
Forsberg, Emil.
Lindström, Viktor.

March 7, 2024

Introduction to the problem

The main force exerted between galaxies is the gravitational force. Newton's law of gravity states that in two dimensions, the gravitational force exerted on body i from j is given by the following equation:

$$F_{ij} = -\frac{Gm_i m_j}{r_{ij}^3} \mathbf{r}_{ij} \quad (1)$$

Here G is the gravitational constant, m_i and m_j are the masses for body i and j respectively, $r_{ij} = \sqrt{(x_i - x_j)^2 + (y_i - y_j)^2}$ is the distance between the bodies, and $\mathbf{r}_{ij} = (x_i - x_j)\hat{e}_x + (y_i - y_j)\hat{e}_y$. To deal with instability problems we modify Newton's law of gravity to include Plummer spheres:

$$F_{ij} = -\frac{Gm_i m_j}{(r_{ij} + \epsilon)^3} \mathbf{r}_{ij} \quad (2)$$

This is an N-body problem, the total gravitational force (exerted by all other bodies) for each body will be calculated to update the velocity and position over time. This can be done with, for example, the symplectic Euler method:

$$\begin{aligned} a_i^n &= \frac{F_i^n}{m_i} \\ u_i^{n+1} &= u_i^n + \Delta t a_i^n \\ x_i^{n+1} &= x_i^n + \Delta t u_i^{n+1} \end{aligned}$$

For this simulation the following values were used: $\epsilon_0 = 10^{-3}$, $G = 100/N$, and $\Delta t = 10^{-5}$.

Solution

It was chosen to use a struct to store the position, velocity, mass, and brightness of each particle. The code reads all the data from the input file and stores it in N particle structures. This is a very efficient way of storing the current values for these parameters, another way would be to store everything in a matrix but this is not very efficient. The code then sends all of these particles as input in the step function. The step function calculates all the forces and updates the velocity and position accordingly for all particles. Memory is allocated to the local variables F_x and F_y with `malloc` to store the forces and freed when the function is done, and it is no longer needed to avoid memory leaks. Step is called `nstep` amount of times as a loop in the main function. It is also possible to split the step function into smaller functions, for example, one function that calculates the force, and one that steps and uses symplectic Euler. But this introduces more function calls and "jumps" in memory. We also realized that the `O3` flag was not doing proper vectorization since the code was updating the i :th element of F_x and F_y in the inner loop. This issue was fixed with a temporary variable.

Discussion of performance

Optimization techniques were used regularly throughout the development of this project. This means that in table 1, "None" optimization actually means some optimization. For example, when calculating the force in each time step, two nested loops are needed. The value $m_i \cdot G$ is constant in the inner loop, so the value is calculated once in the outer one. Actually, producing constant values was done regularly. ϵ_0 was defined at the top of the program. Other constants, such as Δt and G , depended on the input of the user. These were set as constant in main.

Some optimizations were done in hindsight. The final best time can be seen under "manual optimization" in table 1. When calculating the force over all the particles, it was deemed smarter to use two separate loops over j to i and $i + 1$ to N in the inner loop so an if statement could be avoided. (Since it should be summed when $i \neq j$). When using this scheme, the Arrhenius machine completed the $N = 3000$ file after 33 seconds. However, we came up with a strategy to just use one loop instead, going from $j = i + 1$ to $j < N$ in the inner loop. This produced the same result but completed it in 19.7 seconds instead, a pretty significant improvement. Another thing that could be optimized was the initialization of the force vectors in each time step. Firstly, this was done with a separate loop. Afterwards, it was tried to do it in the inner loop so the separate loop could be disposed of. This produced the correct result, but was ultimately slower. This is likely because it tried to reach the value of the vector N^2 times, instead of N times as with the separate loop. It was deemed impossible to do this N times without producing a separate loop. Another possible optimization was to allocate the temporary force arrays on the stack instead of the heap in the step function. However, this did virtually nothing. This is strange, since static memory is generally faster than dynamic.

As previously mentioned the simulation was ran on Arrhenius, which has an Intel(R) Xeon(R) CPU E5520 @2.27GHz with gcc version 11.4.0.

The simulation-time was measured using the linux command "time" in the terminal, with the results for the different optimization methods being presented in table 1 below. Here, 3000 particles were simulated over 100 time-steps.

Table 1: Measured time for simulations with different optimisations, N=3000, nstep=100

Optimisation method	Time for simulation [s]
None	33.0
Manual optimisation	19.7
Manual optimisation with -O3	10.5
Manual optimisation with -O2	10.5
Manual optimisation with -Ofast	10.5

As seen in table 1 above, our original script (referred to as None) took 32.9 seconds to finish.

After reviewing the script, after it's completion, and making the optimization-changes mentioned above, the simulation time was reduced to 19.7 seconds.

Some optimization flags (-O3, -O2, -Ofast) were also implemented on the manually optimized version. this was done by using the flag "-O3", "-O2", and "-Ofast" respectively in the terminal. The resulting time using the respective flags shows a clear decrease in

time, but there seems to be no difference between the three optimization-flags.

Lastly, the time complexity of the program was studied. Eight different simulations were done with different N and 100 timesteps. The results can be seen in figure 1.

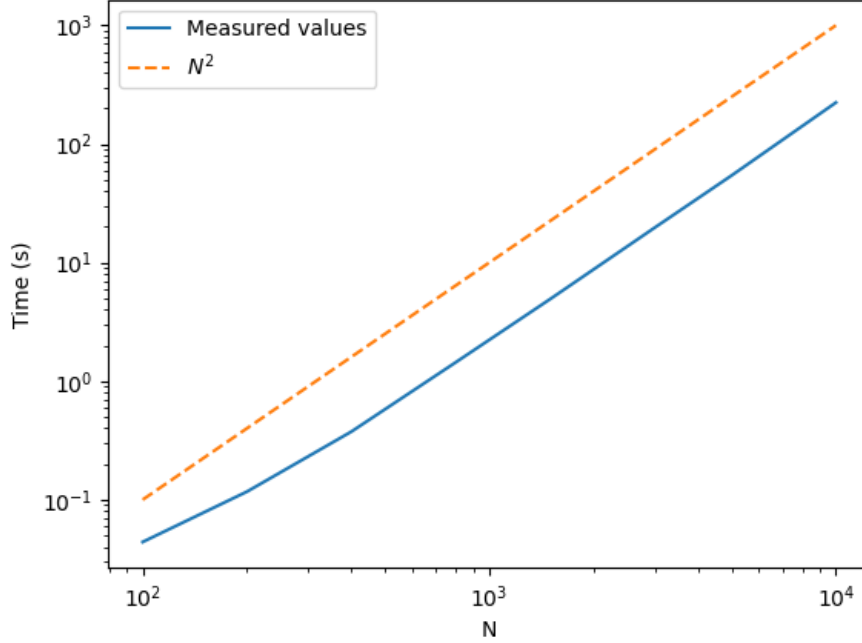


Figure 1: The measured and expected times in a loglog-plot.

When applying linear regression to the logarithmic values, the slope was ≈ 1.9 , which is pretty close to 2. Table 2 shows the values explicitly.

Table 2: Values of N and t

N	$t(s)$
100	0.044
200	0.117
400	0.375
800	1.44
1500	4.93
3000	19.69
5000	54.22
10000	223.5

Here it clearly shows that the time complexity grows as $\mathcal{O}(N^2)$.